



Uber

- ① Case studies → 4 steps
- ② Redis
- ③ Anshuman's Table DB to use
- ④ Industry - Anshuman
+ Transactions

- ① Features
- ② Estimation of Scale
- ③ Design TradeOffs
- ④ Deepdive

① Features

- ① Request / Book cabs
- ② Drivers — get rides
- ③ Matching → 😊
- ④ ⋮

② Estimation of scale

0.1 M drivers / big city

$$0.1 M \times 100 = \underline{10 M} \text{ drivers / globe}$$

← 20M drivers 😊 😊

$$\left[\begin{array}{ll} 7B \times 2\% & \Rightarrow \underline{140M} \text{ users} \times \text{Uber} \\ 7000M \times 2 & = 140M \end{array} \right.$$

③

Design TradeOffs

① low latency - matching problems




②

⑤

Design Deepdive

Quad Trees

① Recompute step — build a quad Tree



②

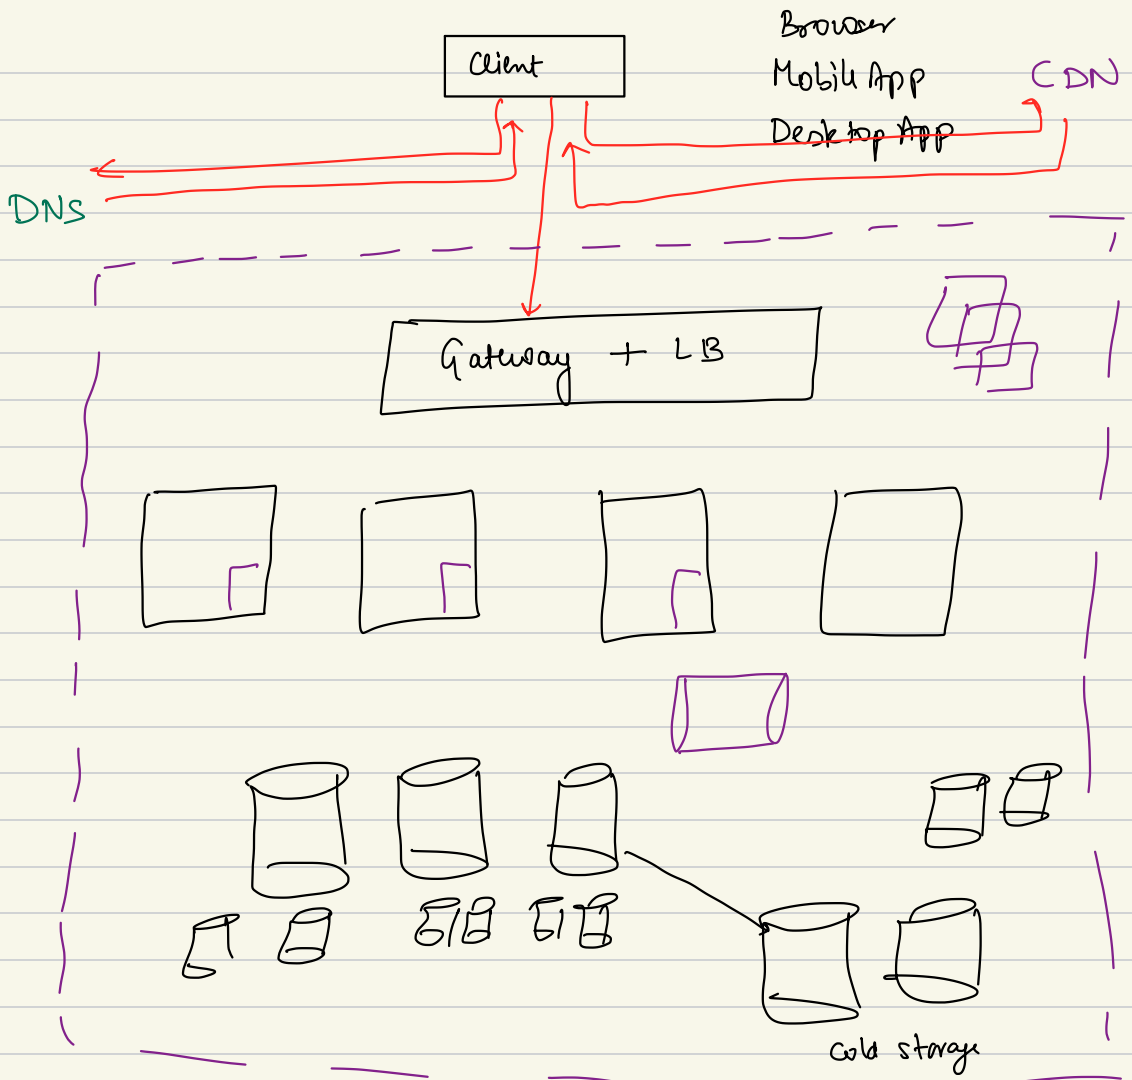
query \rightarrow Find Nearby Places

[piggyback on Quad Tree
and answer]

③

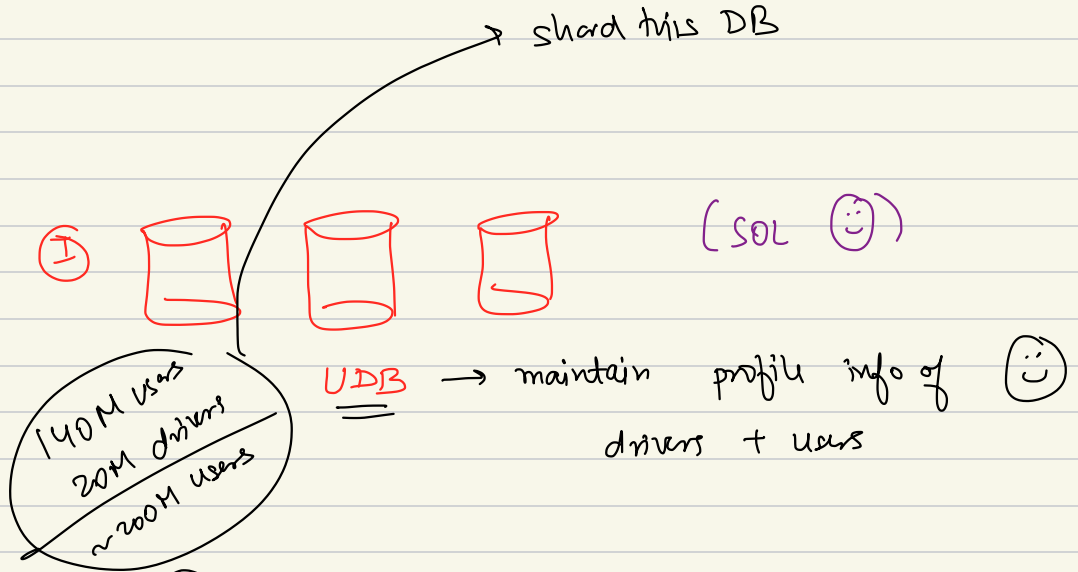
places — added

— removed

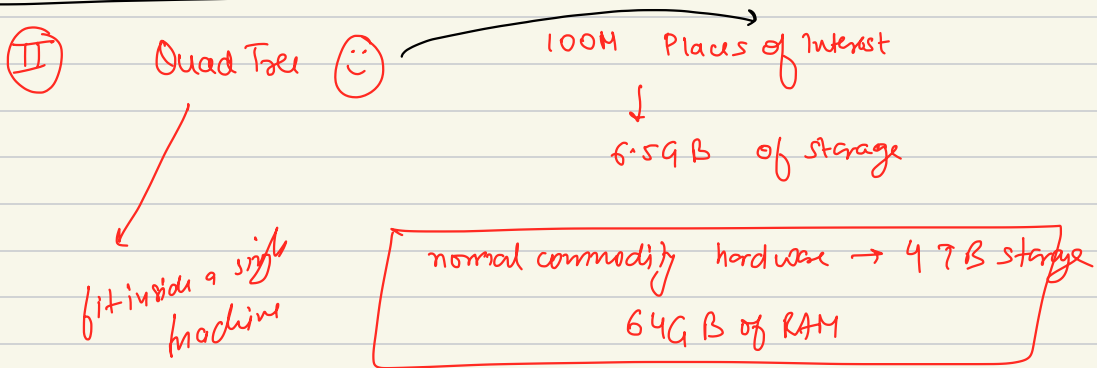


Uber


Federation of DB




Database Federation 😊 😊



Uber Region - Delhi NCR




Uber Region - Jaipur




+ Netmore
+ Alwa

Uber Region - Paris



Intra City






③

Trips DB

SOL



J10


④

Detailed Location Trackers DB



wide column DB

Uber - as a collection of
multiple Regional Uber systems



zoom into any 1 Uber Region





→ 100M → 6.5 GB ≈
0.1M → few MBs 😊 😊

d-id → last location; last grid-cell-id
Global Cache


driver's loc

Gateway + LB

Quad Tree



UDB




Trips



10:39 - 10:49 PM


+ entry ①



Location DB

Step 1: Bootstrap your quad Tree

[some reference data to
approximate]







Driver

periodically

l_{min} → lat, long

Global Cache ≡ Redis Cluster




I

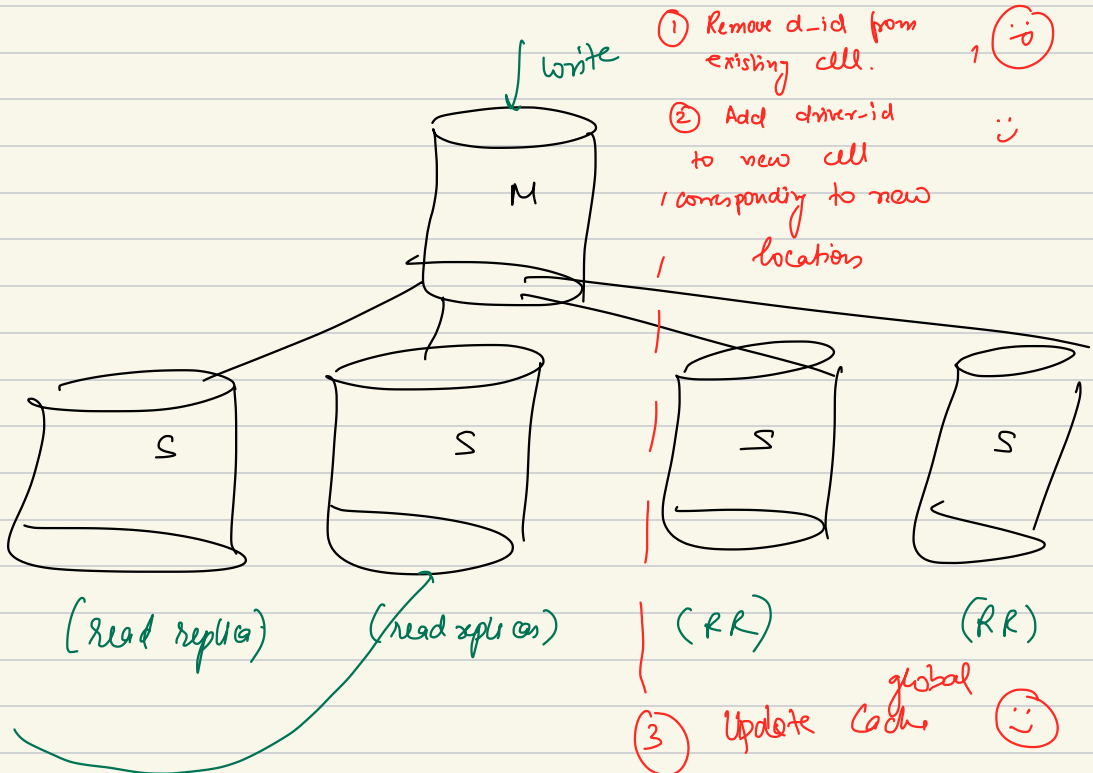
d-id \rightarrow Last (x,y) location \Leftarrow
Last cell-id

II

cell-id \rightarrow



Quad Tree

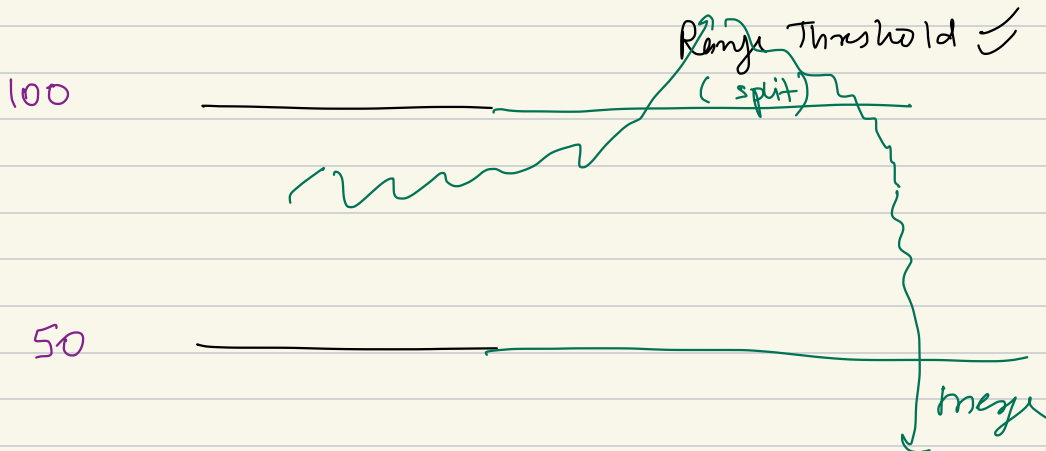


So, Uber system design seems like a clever use case
of Quad Trees

- ① We have diff. subsystems for every region.
- ② Inside each subsystem, we have an independent Quad Tree.
- ③ We are cleverly using Redis global cache clients and/or slave (Read Replicas) of the Quad Tree machine to minimize the write requests to the Quad Tree.


PROBLEM OF CONTINUOUS MERGING
and SPLITTING

1. Point Threshold X

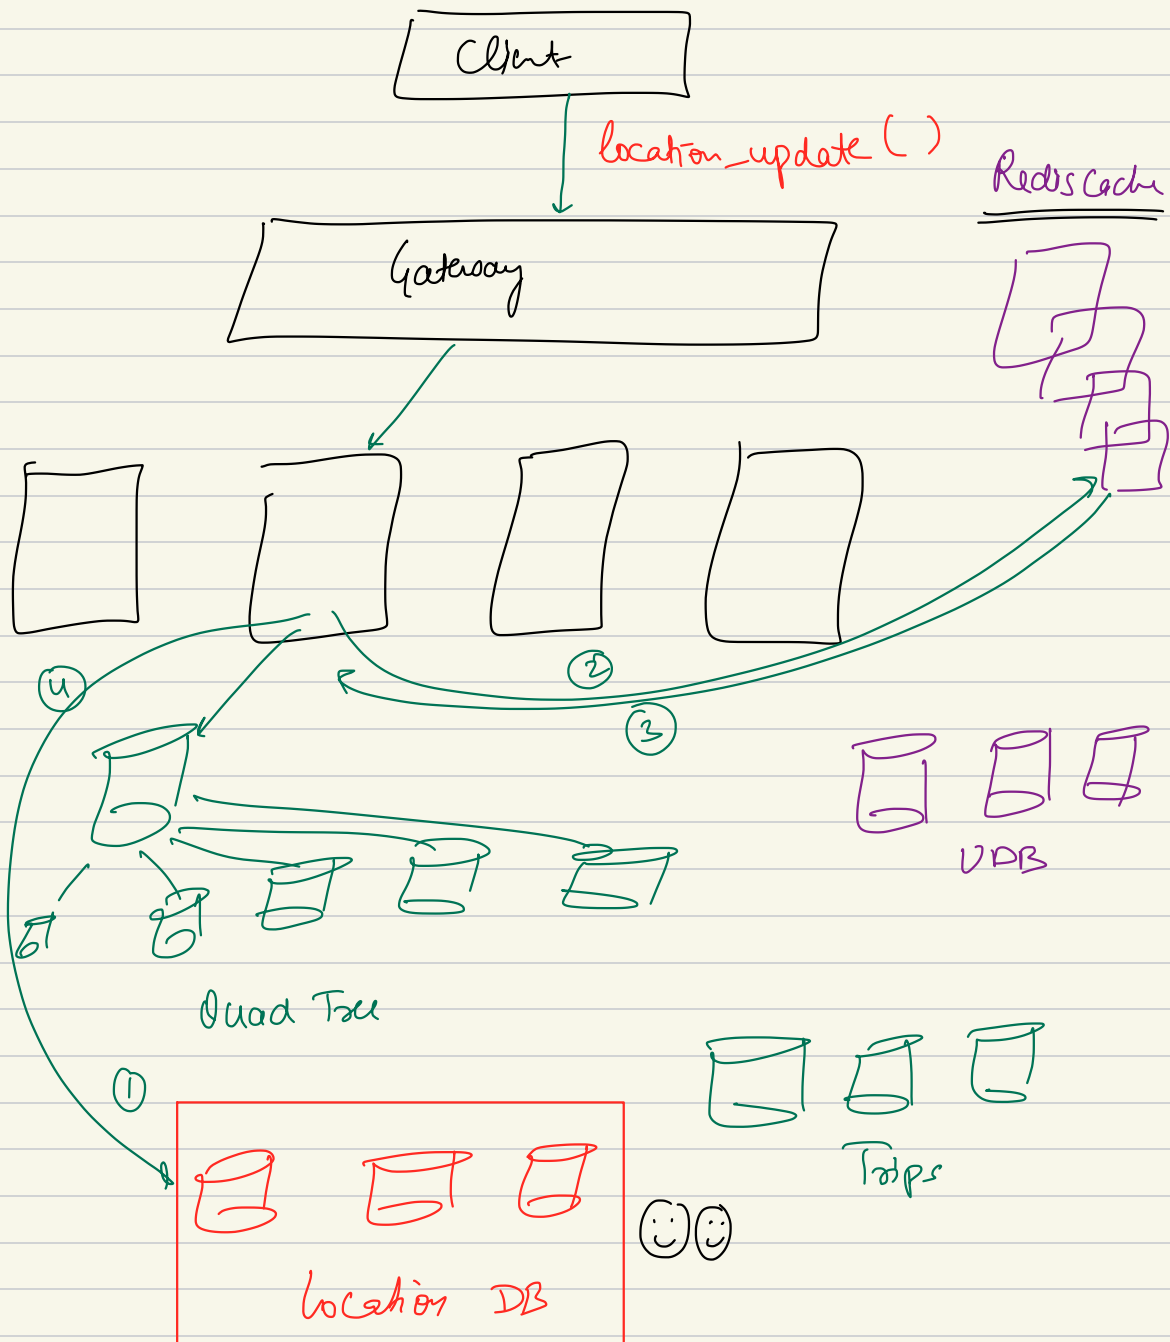


2


Uber decides that grid cells
are only going to change once
ever 1 hour / 2 hour



Driver Location Flow 😊




Ride Booking Flow



Active Drivers

(websocket)

HTTP Long Polling



do you have a ride for me?

Server

